

Install necessary packages

```
In [1]: # !pip install tensorflow keras pillow numpy tqdm
```

import all the necessary packages

```
In [2]: import string
import numpy as np
from PIL import Image
import os
from pickle import dump, load
import numpy as np

from keras.applications.xception import Xception, preprocess_input
from keras.preprocessing.image import load_img, img_to_array
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.utils import to_categorical
# from keras.utils import to_categorical
from keras.layers.merge import add
from keras.models import Model, load_model
from keras.layers import Input, Dense, LSTM, Embedding, Dropout

# small library for seeing the progress of loops.
from tqdm import tqdm_notebook as tqdm
tqdm().pandas()
```

C:\Users\Client\AppData\Local\Temp\ipykernel_6484\379097116.py:20: TqdmDeprecationWarning: This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
tqdm().pandas()

Getting and performing data cleaning

```
In [3]: # Loading a text file into memory
def load_doc(filename):
    # Opening the file as read only
    file = open(filename, 'r')
    text = file.read()
    file.close()
    return text

# get all imgs with their captions
def all_img_captions(filename):
    file = load_doc(filename)
    captions = file.split('\n')
    descriptions = {}
    for caption in captions[:-1]:
        img, caption = caption.split('\t')
        if img[:-2] not in descriptions:
            descriptions[img[:-2]] = [caption]
        else:
            descriptions[img[:-2]].append(caption)
    return descriptions

#Data cleaning- Lower casing, removing punctuations and words containing numbers
def cleaning_text(captions):
    table = str.maketrans('', '', string.punctuation)
    for img, caps in captions.items():
        for i, img_caption in enumerate(caps):
```

```

img_caption.replace("-", "")
desc = img_caption.split()

#converts to Lowercase
desc = [word.lower() for word in desc]
#remove punctuation from each token
desc = [word.translate(table) for word in desc]
#remove hanging 's and a
desc = [word for word in desc if(len(word)>1)]
#remove tokens with numbers in them
desc = [word for word in desc if(word.isalpha())]
#convert back to string

img_caption = ' '.join(desc)
captions[img][i]= img_caption
return captions

def text_vocabulary(descriptions):
    # build vocabulary of all unique words
    vocab = set()

    for key in descriptions.keys():
        [vocab.update(d.split()) for d in descriptions[key]]

    return vocab

#All descriptions in one file
def save_descriptions(descriptions, filename):
    lines = list()
    for key, desc_list in descriptions.items():
        for desc in desc_list:
            lines.append(key + '\t' + desc )
    data = "\n".join(lines)
    file = open(filename,"w")
    file.write(data)
    file.close()

```

In [4]:

```

# Set these path according to project folder in you system
dataset_text = "C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/Flickr8k_text"
dataset_images = "C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/Flicker8k_Data:

#we prepare our text data
filename = dataset_text + "/" + "Flickr8k.token.txt"
#Loading the file that contains all data
#mapping them into descriptions dictionary img to 5 captions
descriptions = all_img_captions(filename)
print("Length of descriptions =" , len(descriptions))

#cleaning the descriptions
clean_descriptions = cleaning_text(descriptions)

#building vocabulary
vocabulary = text_vocabulary(clean_descriptions)
print("Length of vocabulary = ", len(vocabulary))

#saving each description to file
save_descriptions(clean_descriptions, "C:/Users/Client/Downloads/Image-Caption-Generator Fl.

```

Length of descriptions = 8092
Length of vocabulary = 8763

Extracting the feature vector from all images (TAKES TIME)

In [5]:

```

def extract_features(directory):

```

```

model = Xception( include_top=False, pooling='avg' )
features = {}
for img in tqdm(os.listdir(directory)):
    filename = directory + "/" + img
    image = Image.open(filename)
    image = image.resize((299,299))
    image = np.expand_dims(image, axis=0)
    #image = preprocess_input(image)
    image = image/127.5
    image = image - 1.0

    feature = model.predict(image)
    features[img] = feature
return features

#2048 feature vector
features = extract_features(dataset_images)
dump(features, open("C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/features.p"

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kernels_notop.h5
83689472/83683744 [=====] - 31s 0us/step
83697664/83683744 [=====] - 31s 0us/step
C:\Users\Client\AppData\Local\Temp\ipykernel_6484\3734661667.py:4: TqdmDeprecationWarning: This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
for img in tqdm(os.listdir(directory)):

In [6]: features = load(open("C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/features.p"

Loading dataset for Training the model

In [7]:

```

#Load the data
def load_photos(filename):
    file = load_doc(filename)
    photos = file.split("\n")[:-1]
    return photos

def load_clean_descriptions(filename, photos):
    #loading clean descriptions
    file = load_doc(filename)
    descriptions = {}
    for line in file.split("\n"):

        words = line.split()
        if len(words)<1 :
            continue

        image, image_caption = words[0], words[1:]

        if image in photos:
            if image not in descriptions:
                descriptions[image] = []
            desc = '<start> ' + " ".join(image_caption) + ' <end>'
            descriptions[image].append(desc)

    return descriptions

def load_features(photos):
    #loading all features
    all_features = load(open("C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/fe
    #selecting only needed features
    features = {k:all_features[k] for k in photos}

```

```

        return features

filename = dataset_text + "/" + "Flickr_8k.trainImages.txt"

#train = loading_data(filename)
train_imgs = load_photos(filename)
train_descriptions = load_clean_descriptions("C:/Users/Client/Downloads/Image-Caption-Generator/Flickr8k/trainDescriptions.txt")
train_features = load_features(train_imgs)

```

Tokenizing the vocabulary

```

In [8]: #converting dictionary to clean list of descriptions
def dict_to_list(descriptions):
    all_desc = []
    for key in descriptions.keys():
        [all_desc.append(d) for d in descriptions[key]]
    return all_desc

#creating tokenizer class
#this will vectorise text corpus
#each integer will represent token in dictionary

from keras.preprocessing.text import Tokenizer

def create_tokenizer(descriptions):
    desc_list = dict_to_list(descriptions)
    tokenizer = Tokenizer()
    tokenizer.fit_on_texts(desc_list)
    return tokenizer

# give each word an index, and store that into tokenizer.p pickle file
tokenizer = create_tokenizer(train_descriptions)
dump(tokenizer, open('C:/Users/Client/Downloads/Image-Caption-Generator/Flickr8k/tokenizer.pkl', 'wb'))
vocab_size = len(tokenizer.word_index) + 1
vocab_size

```

Out[8]: 7577

```

In [9]: #calculate maximum length of descriptions
def max_length(descriptions):
    desc_list = dict_to_list(descriptions)
    return max(len(d.split()) for d in desc_list)

max_length = max_length(descriptions)
max_length

```

Out[9]: 32

Create Data generator

```

In [10]: #create input-output sequence pairs from the image description.

#data generator, used by model.fit_generator()
def data_generator(descriptions, features, tokenizer, max_length):
    while 1:
        for key, description_list in descriptions.items():
            #retrieve photo features
            feature = features[key][0]
            input_image, input_sequence, output_word = create_sequences(tokenizer, max_length, description_list, feature)
            yield ([input_image, input_sequence], output_word)

```

```

def create_sequences(tokenizer, max_length, desc_list, feature):
    X1, X2, y = list(), list(), list()
    # walk through each description for the image
    for desc in desc_list:
        # encode the sequence
        seq = tokenizer.texts_to_sequences([desc])[0]
        # split one sequence into multiple X,y pairs
        for i in range(1, len(seq)):
            # split into input and output pair
            in_seq, out_seq = seq[:i], seq[i]
            # pad input sequence
            in_seq = pad_sequences([in_seq], maxlen=max_length)[0]
            # encode output sequence
            out_seq = to_categorical([out_seq], num_classes=vocab_size)[0]
            # store
            X1.append(feature)
            X2.append(in_seq)
            y.append(out_seq)
    return np.array(X1), np.array(X2), np.array(y)

#You can check the shape of the input and output for your model
[a,b],c = next(data_generator(train_descriptions, features, tokenizer, max_length))
a.shape, b.shape, c.shape

```

Out[10]: ((47, 2048), (47, 32), (47, 7577))

Defining the CNN-RNN model

```

In [11]: from tensorflow.keras.utils import plot_model

# define the captioning model
def define_model(vocab_size, max_length):

    # features from the CNN model squeezed from 2048 to 256 nodes
    inputs1 = Input(shape=(2048,))
    fe1 = Dropout(0.5)(inputs1)
    fe2 = Dense(256, activation='relu')(fe1)

    # LSTM sequence model
    inputs2 = Input(shape=(max_length,))
    se1 = Embedding(vocab_size, 256, mask_zero=True)(inputs2)
    se2 = Dropout(0.5)(se1)
    se3 = LSTM(256)(se2)

    # Merging both models
    decoder1 = add([fe2, se3])
    decoder2 = Dense(256, activation='relu')(decoder1)
    outputs = Dense(vocab_size, activation='softmax')(decoder2)

    # tie it together [image, seq] [word]
    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
    model.compile(loss='categorical_crossentropy', optimizer='adam')

    # summarize model
    print(model.summary())
    plot_model(model, to_file='C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/m

    return model

```

TRAINING SECTION

```

In [12]: # train our model
print('Dataset: ', len(train_imgs))
print('Descriptions: train=', len(train_descriptions))

```

```

print('Photos: train=', len(train_features))
print('Vocabulary Size:', vocab_size)
print('Description Length: ', max_length)

model = define_model(vocab_size, max_length)
print(model, 'model')
epochs = 10
steps = len(train_descriptions)
# making a directory models to save our models
os.mkdir("C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/models")
for i in range(epochs):
    generator = data_generator(train_descriptions, train_features, tokenizer, max_length)
    model.fit_generator(generator, epochs=1, steps_per_epoch=steps, verbose=1)
    model.save("C:/Users/Client/Downloads/Image-Caption-Generator Flickr8k/models/model_" +

```

Dataset: 6000
 Descriptions: train= 6000
 Photos: train= 6000
 Vocabulary Size: 7577
 Description Length: 32
 Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_3 (InputLayer)	[(None, 32)]	0	[]
input_2 (InputLayer)	[(None, 2048)]	0	[]
embedding (Embedding)	(None, 32, 256)	1939712	['input_3[0][0]']
dropout (Dropout)	(None, 2048)	0	['input_2[0][0]']
dropout_1 (Dropout)	(None, 32, 256)	0	['embedding[0][0]']
dense (Dense)	(None, 256)	524544	['dropout[0][0]']
lstm (LSTM)	(None, 256)	525312	['dropout_1[0][0]']
add_12 (Add)	(None, 256)	0	['dense[0][0]', 'lstm[0][0]']
dense_1 (Dense)	(None, 256)	65792	['add_12[0][0]']
dense_2 (Dense)	(None, 7577)	1947289	['dense_1[0][0]']

=====
 =====
 Total params: 5,002,649
 Trainable params: 5,002,649
 Non-trainable params: 0

None

('You must install pydot (`pip install pydot`) and install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) ', 'for plot_model/model_to_dot to work.')
 <keras.engine.functional.Functional object at 0x000001FF70EEDA00> model

C:\Users\Client\AppData\Local\Temp\ipykernel_6484\409286170.py:16: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

model.fit_generator(generator, epochs=1, steps_per_epoch=steps, verbose=1)
 6000/6000 [=====] - 2255s 375ms/step - loss: 4.5057

C:\Users\Client\anaconda3\envs\ds\lib\site-packages\keras\engine\functional.py:1410: CustomMaskWarning: Custom mask layers require a config and must override get_config. When loading, the custom mask layer must be passed to the custom_objects argument.

layer_config = serialize_layer_fn(layer)

6000/6000	[=====]	-	2239s	373ms/step	-	loss: 3.6630
6000/6000	[=====]	-	2247s	375ms/step	-	loss: 3.3745
6000/6000	[=====]	-	2280s	380ms/step	-	loss: 3.2001
6000/6000	[=====]	-	2281s	380ms/step	-	loss: 3.0849
6000/6000	[=====]	-	2263s	377ms/step	-	loss: 2.9947
6000/6000	[=====]	-	2285s	381ms/step	-	loss: 2.9279
6000/6000	[=====]	-	2291s	382ms/step	-	loss: 2.8658
6000/6000	[=====]	-	2277s	379ms/step	-	loss: 2.8246
6000/6000	[=====]	-	2267s	378ms/step	-	loss: 2.7829